

Governance and grounding layer for trusted AI content

Summary

AURORA addresses a practical enterprise AI gap: employees have access to AI tools, but the tools are not yet integrated into governed content workflows. The final presentation shows strong willingness to use AI: 74 percent of surveyed employees use AI daily and only 5 percent find the tools hard to use. The barrier is not usability. The barrier is clarity, role-specific guidance, trust, source control, and workflow integration. Only 12 percent see clear expectations, 62 percent miss role examples, and 66 percent would use AI more with better guidance.

Summary

AURORA addresses a practical enterprise AI gap: employees have access to AI tools, but the tools are not yet integrated into governed content workflows. The final presentation shows strong willingness to use AI: 74 percent of surveyed employees use AI daily and only 5 percent find the tools hard to use. The barrier is not usability. The barrier is clarity, role-specific guidance, trust, source control, and workflow integration. Only 12 percent see clear expectations, 62 percent miss role examples, and 66 percent would use AI more with better guidance.

AURORA is the proposed answer for ABN AMRO content workflows. It is not another general chatbot. It is a governed grounding and evaluation layer that turns a vague request into a specific, on-brand, source-backed draft with citations, KPI checks, and a visible audit trail. The proof of concept focuses on ABN AMRO Insights articles in Technology, Media and Telecom because that workflow has real value, approved public source material, clear writing standards, and measurable quality criteria.

The current proof of concept is credible. It includes a seven-stage grounded pipeline, REST API and MCP server, reference agent wrapper, retrieval over 59 approved TMT Insight articles, English and Dutch retrieval assets, six editorial and domain profiles, citation checking, typed outputs, per-stage reasoning, and deterministic fallback paths. The evaluation layer is based on a 135-entry KPI catalogue, with active draft checks across citation validity, traceability, approved-source status, factuality, truthfulness, relevance, privacy, groundedness, and comprehensiveness.

The recommendation is to approve AURORA for a controlled pilot and production-hardening track. It should not be rolled out broadly until access control, audit storage, EU-resident deployment, monitoring, incident handling, and operating ownership are in place.

Why this matters

The current content workflow contains a costly rewrite loop. A domain expert or analyst drafts content, a content specialist edits and flags issues, the draft returns for rewriting, and the process repeats until it meets quality, tone, source, and compliance expectations.

AURORA changes that pattern. The aim is not to remove human judgement. The aim is to make the first draft closer to usable by embedding approved sources, editorial standards, role profiles, prompt refinement, citations, and evaluation into the workflow from the start. Domain experts can focus on substance. Content specialists can focus on final judgement instead of preventable rework.

The business case should be measured by workflow effectiveness, not token use. The presentation estimates 20 to 40 hours saved each week across a 42-person content team if integration deepens. A realistic pilot target is to move lower-savings users up by 2 to 3 hours per person per week. This is plausible because the survey shows demand exists; employees mainly need clearer guidance and better integration.

The main value drivers are:

- Faster throughput through fewer rewrite loops.
- Better use of domain experts and content specialists.
- Stronger source traceability and citation discipline.
- Lower duplication through one reusable governance layer.

- Lower AI risk through approved sources, checks, blocking rules, and human review.
- Lower vendor dependence because ABN AMRO's standards stay outside the model provider.

How AURORA works

AURORA uses a seven-stage pipeline.

1. Intent classification detects task type, sector, topic keywords, and output language.
2. Profile selection chooses workflow and domain expert profiles, such as drafter, reviewer, curator, TMT generalist, cybersecurity specialist, or media specialist.
3. Context retrieval searches approved corpora before writing, including English and Dutch article corpora, writing guides, and Insights style rules.
4. Prompt refinement turns the raw request into a clearer, complete, source-aware instruction.
5. Conditional rerun repeats earlier steps if clarification changes the task.
6. Content generation produces a grounded draft using the refined prompt, selected profiles, retrieved snippets, and citation requirements.
7. Evaluation checks the output against active KPI checks and identifies whether the result is usable, blocked with reasons, or pending human sign-off.

Strategically, AURORA should evolve from a fixed pipeline into a modular enterprise tool layer. Approved AI agents should be able to call only the capability they need: retrieval, prompt refinement, citation checking, evaluation, translation support, or the full pipeline.

Current maturity

Built now:

- Live grounded pipeline over approved TMT Insight articles.
- REST API and MCP server.
- Reference agent wrapper.
- Retrieval over 59 approved articles.
- English and Dutch support.
- Six editorial and domain profiles.
- Citation checking and per-stage reasoning.
- Deterministic fallback paths.

Partially implemented:

- Three-tier KPI evaluation model.
- Local OpenAI Agent integration.
- Model-provider abstraction.
- Complete audit functionality.
- Enterprise AI and Copilot-style integration architecture.

Required before production:

- Entra ID authentication and role-based access.
- Durable and immutable audit database.
- EU-resident deployment.

- Production-grade hybrid retrieval backend.
- Broader evaluator coverage.
- Monitoring, incident management, and operating ownership.

Next Step

AURORA should move forward with disciplined scope. The proof of concept demonstrates a strong pattern for governed AI adoption: approved sources before writing, role-specific guidance during generation, KPI evaluation after drafting, and human sign-off where needed.

The next step is a controlled pilot, not broad deployment. If the pilot confirms fewer rewrite loops, better first-pass quality, stronger traceability, and improved user confidence, AURORA can become a reusable governance layer for trusted AI content across more ABN AMRO workflows.